

Exercise 2.1

1. The sequential version took 5.6 seconds.
2. The time taken by 10-thread version is 1.4 seconds which is really fast than the sequential one.
3. No, it produces different result each time I run the program.
4. No, the synchronized is not required here because the `get()` method is not accessed by any threads and the count variable state is not changed.

Exercise 2.2

1. The program takes 5.9 seconds to execute.
2. `MyAtomicInteger` is implemented in `TestCountFactors.java`.
3. Yes, it is displaying the same answer and the time to execute is 1.5 seconds.
4. No, we cannot implement `MyAtomicInteger` without synchronization because volatile field can only guarantee visibility but not atomicity. In absence of synchronization mutual exclusion cannot be obtained and thread will see stale data.
5. Yes, it is important to declare `AtomicInteger` as final so that its state is immutable.

Exercise 2.3

1. Cache field is declared volatile because when a thread sets the volatile cache field to reference a new `OneValueCache`, the new cached data becomes immediately visible to other threads. [answer from page 33 from the book]
2. Both need to be declared final because when a thread acquires a reference to immutable object it need never worry about another thread modifying its state.

Exercise 2.4

1. Implemented histogram2 in `SimpleHistogram.java`
 - a. `increment()` and `getCount()` should be made synchronized in order to achieve atomicity.
 - b. The `getSpan()` method does not need to be synchronized because it is not the part of atomic operation.
2. Implemented in `SimpleHistogram.java`
3. Yes, we can remove synchronized now because each integer is atomic now and no thread can interfere in between the operation making it atomic without synchronized keyword. The results are correctly printed.
4. The results are correctly printed.
5. Implemented in `SimpleHistogram.java`
6. Implemented in `SimpleHistogram.java`

Exercise 2.5

1. /

2.

Memoizer1 gives us no poor concurrency. No parallelism at all actually.

Output:

```
class Memoizer1
115000
```

```
real    0m26.428s
user    0m24.228s
sys     0m1.800s
```

3.

Memoizer2 gives us a decent degree of parallelism but there are many cache misses (because there is still a window of vulnerability). $132486 - 115000 = 17486$

cache misses. The problem here is that two or more threads can start the computation at the same time for the same number.

```
class Memoizer2
132486
```

```
real    0m19.790s
user    0m37.720s
sys     0m0.088s
```

4.

Memoizer3 is faster than all the previous executions but there are still some cache misses. This happens because the if statement in the compute method is non-atomic. The Figure 5.4 page68 (Goetz) depicts this problem very well.

```
class Memoizer3
115092
```

```
real    0m16.857s
user    0m25.260s
sys     0m0.744s
```

5.

As mentioned in the code, a hybrid of Memoizer3 and Memoizer. We can see that from the degree of parallelism and speed of execution, they give roughly the same results. But, in this case we also get no cache misses which is because now we use an atomic putIfAbsent call.

```
class Memoizer4
115000
```

```
real    0m16.754s
user    0m25.508s
sys     0m0.744s
```

6.

No cache misses but slower than Memoizer4.

```
class Memoizer5
115000
```

```
real    0m23.842s
user    0m38.272s
sys     0m1.032s
```

7.

This simple implementation actually gives good results in the degree of parallelism, cache hits and speed of execution.

```
class Memoizer0
115000
```

```
real    0m18.424s
user    0m30.040s
sys     0m1.012s
```

Exercise 2.6

1. The thread that finishes last does not produce the same number as the main thread.

Example output:

main

main finished 40000000

fresh 1 stops: 23909808

fresh 0 stops: 24015599

2. The problem here is that we have synchronized on a Boxed primitive (a Long in this case). Boxed primitives are immutable which means that when we do "count++" this actually ends up being executed as smth like this:
`count = new Long(count + 1);`

So basically, we end up synchronizing on a different reference (Object) each time we call count++ which is not how synchronization properly works.

3.

I solved this by creating an explicit lock:
`final static Object lock = new Object();`

main

main finished 40000000

fresh 1 stops: 39670696

fresh 0 stops: 40000000